

NAINIKA DWIVEDI

Software Engineer | AI/GenAI & LLM Systems

+91-9555972430 | nainikadwivedi526@gmail.com | Noida, India | linkedin.com/in/nainika-dwivedi-6b986a1a6 | 5+ Years Experience

SUMMARY

Software Engineer with 5+ years of experience building production-oriented LLM applications, RAG pipelines, and agentic AI workflows in Python. Hands-on experience integrating and benchmarking multiple LLMs (OpenAI GPT, Google Gemini) via LangChain and LangGraph, designing automated evaluation frameworks for NL-to-SQL generation across enterprise data warehouses (Amazon Redshift, Databricks, Snowflake), and shipping FastAPI-based GenAI services. Background in embedded C++ systems programming brings strong fundamentals in performance, concurrency, and production reliability to AI engineering work.

KEY SKILLS

Programming Languages: Python, C++, C

GenAI / LLM: LLM APIs (OpenAI GPT, Google Gemini), LangChain, LangGraph, Agentic AI, Retrieval-Augmented Generation (RAG), Prompt Engineering, Vertex AI, Hugging Face, LLM Evaluation & Benchmarking

Backend / APIs: FastAPI, REST APIs, Microservices Integration

Data & Cloud: GCP, BigQuery, Amazon Redshift, Snowflake, Databricks, PySpark, Pandas, NumPy, TensorFlow, PyTorch

Databases: SQL, MySQL, Oracle DBMS, RDBMS

Systems / Other: Embedded Programming, Linux, Git

WORK EXPERIENCE

Software Engineer — GenAI/LLM Systems | GlobalLogic

May 2023 – Jul 2026

- Designed and led an automated LLM evaluation pipeline benchmarking OpenAI GPT and Google Gemini (multiple Flash variants) on NL-to-SQL generation, covering 500+ queries across 30+ datasets.
- Evaluated model-generated SQL across three enterprise dialects — Amazon Redshift, Databricks, and Snowflake — scoring accuracy, latency, execution time, and Valid Efficiency Score (VES) to guide model selection.
- Built RAG pipelines for enterprise knowledge retrieval, implementing document ingestion, chunking, embeddings, semantic search, and vector database integration to ground LLM responses in relevant context.
- Developed agentic AI workflows using LangChain and LangGraph, orchestrating retrieval, reasoning, tool execution, and LLM-based decision-making.
- Integrated Google Vertex AI and Gemini models into production GenAI workflows; developed prompt templates and structured-output pipelines for consistent, evaluable responses.
- Shipped API-based GenAI services using FastAPI, exposing LLM, RAG, and agentic workflows to downstream applications.
- Collaborated within cross-functional AI engineering teams that scaled from an initial 25–30 members down to a focused 6–8 engineer core team, maintaining delivery velocity through the transition.
- Applied prompt engineering and retrieval optimization to iteratively improve response quality, relevance, and reliability across evaluation cycles.

Software Engineer — Embedded Systems (C++) | HARMAN Connected Services

Sep 2021 – Feb 2023

- Developed features for OBD and non-OBD automotive dongle devices using C++ and Linux, including VIN validation and vehicle-data functionality driven by sensor/device input.
- Built firmware interaction commands using multithreading and concurrent programming, strengthening systems-level fundamentals now applied to production AI service reliability.
- Worked across Linux, C/C++, Git, and core networking/OS concepts in a hardware-integrated environment.

Trainee Analyst — Python | Capgemini

Mar 2021 – Jun 2021

- Gained applied experience in Python programming and database management systems as part of foundational training.

EDUCATION

B.Tech, Computer Science and Engineering

Noida Institute of Engineering & Technology, Greater Noida

Aug 2017 – Aug 2021

Grade: 8.5/10